

# Socket Performance Analysis Report

---

**Assignment:** GRS\_PA02

**Student ID:** MT25043

**Course:** GRS

**Date:** February 2026

**GitHub:** [https://github.com/shivam697/GRS\\_PA02](https://github.com/shivam697/GRS_PA02)

---

---

## Table of Contents

1. Implementation Overview
  2. Part A: Three Socket Approaches
  3. Part E: Performance Analysis
  4. Experimental Results
  5. AI Usage Declaration
- 
- 

## Implementation Overview

This assignment implements and compares **three socket communication approaches**:

Approach	Method	Key Features
Two-Copy	send() + recv()	Baseline with memcpy to intermediate buffer

One-Copy	sendmsg() + iovec	Scatter-gather I/O, eliminates memcpy
Zero-Copy	MSG_ZEROCOPY	DMA transfer, async completion

### Test Setup:

- Network namespaces (ns1 ↔ ns2) with veth pairs
  - Message sizes: 1KB, 4KB, 16KB, 64KB
  - Thread counts: 1, 2, 4, 8
  - Profiling: perf stat (cycles, cache misses, context switches)
  - Visualization: 4 matplotlib plots
- 
- 

## Part A: Three Socket Approaches

### A1: Two-Copy Implementation (Baseline)

#### How it works:

User buffers (8 fields)  
 → memcpy() to intermediate buffer [COPY #1]  
 → send() to kernel socket buffer [COPY #2]  
 → Network

#### Code snippet (Server):

```
[c]
char* send_buffer = malloc(g_msg_size);
for (int i = 0; i < 8; i++) {
    memcpy(current_pos, msg->field[i], field_size); // COPY #1
    current_pos += field_size;
}
send(client_socket, send_buffer, g_msg_size, 0); // COPY #2
```

**Performance:** Throughput = 43.90 Gbps @ 64KB, 4 threads

[SCREENSHOT: Terminal showing server accepting connection and client throughput/latency]

## A2: One-Copy Implementation

**Question:** You must explicitly demonstrate which copy has been eliminated.

**Answer:** The `memcpy()` to intermediate buffer is eliminated.

## Comparison:

Step	Two-Copy	One-Copy
------	----------	----------

1	memcpy() 8 fields → buffer	~~Eliminated~~
2	send(buffer) → kernel	sendmsg(iovec[8]) → kernel
Copies	2 copies	1 copy

### Code showing elimination:

[C]

```
// ONE-COPY: No intermediate buffer, direct pointers
struct iovec iov[8];
for (int i = 0; i < 8; i++) {
    iov[i].iov_base = msg->field[i]; // Just pointers, NO memcpy
    iov[i].iov_len = field_size;
}
sendmsg(client_socket, &msg_hdr, 0); // Kernel gathers from 8 sources
```

**Result:** One-copy eliminates the memcpy step but adds scatter-gather overhead.

**Performance:** Throughput = 46.89 Gbps @ 64KB (beats two-copy by 6.8%)


---


## A3: Zero-Copy Implementation

**Question:** You must explain kernel behavior using a diagram.

**Answer:**

### Kernel Behavior Diagram:





### Key Steps:

6. `sendmsg(MSG_ZEROCOPY)`: Kernel pins user pages (no copy, just references)
  7. DMA Transfer: NIC reads directly from user memory
  8. Completion: Interrupt → error queue notification
  9. Application: Polls error queue to know when buffer is free

## Why "Zero-Copy"?

- CPU never copies data
  - DMA does the work
  - Trade-off: Async complexity + page pinning overhead

**Performance:** Throughput = 22.91 Gbps @ 64KB (slower due to loopback testing)

[SCREENSHOT: Zero-copy server output with completion draining]

The screenshot shows a terminal window with multiple tabs and panes. The tabs include 'EXPLORER', 'PROBLEMS', 'OUTPUT', 'DEBUG CONSOLE', 'TERMINAL', and 'PORTS'. The 'TERMINAL' tab is active, displaying a series of commands and their outputs related to file conversion and execution.

Key outputs from the terminal:

- Conversion of 'MT25043\_PA02\_Report.docx' to 'hash' (MD5).
- Execution of 'MT25043\_PA02\_Report.docx' (Word document).
- Conversion of 'MT25043\_Part\_A1\_Client.c' to 'hash' (MD5).
- Execution of 'MT25043\_Part\_A1\_Server.c' (C code).
- Conversion of 'MT25043\_Part\_A2\_Client.c' to 'hash' (MD5).
- Execution of 'MT25043\_Part\_A2\_Server.c' (C code).
- Conversion of 'MT25043\_Part\_C\_Server.c' to 'hash' (MD5).
- Execution of 'MT25043\_Part\_C\_Server.c' (C code).
- Conversion of 'MT25043\_Part\_C\_Shell.c' to 'hash' (MD5).
- Execution of 'MT25043\_Part\_D\_Plottting.py' (Python script).
- Execution of 'E\_coe\_copy\_client' (Windows command).
- Execution of 'E\_zero\_copy\_server' (Windows command).
- Execution of 'E\_zero\_copy\_client' (Windows command).
- Execution of 'E\_zero\_copy\_server' (Windows command).

The right side of the interface shows a 'CHAT' pane with a message from the user: "Informed user that REPORT-md conversion was completed". Below it, a note says: "The Word file is already created! I converted it in my previous response."

At the bottom, there are sections for 'How to Open:' (option 1: Open in Microsoft Word) and 'Option 2: In External Application' (using LibreOffice Writer). A note at the bottom right says: "Ready to edit and add your screenshots!"


## Experimental Results

## Part B: Profiling with perf stat

- Integrated perf stat to collect: cycles, instructions, L1/LLC cache misses, branches, context switches
  - CSV output format for automated parsing

## Part C: Automated Experiments

- Bash script runs 48 experiments ( $3$  implementations  $\times$   $4$  thread counts  $\times$   $4$  message sizes)
  - Network namespaces (ns1/ns2) for isolated testing
  - Results saved to MT25043\_Part\_C\_Results.csv






## Part D: Visualization

Generated 4 plots using matplotlib:

10. Throughput vs Message Size - Shows performance scaling
  11. Latency vs Thread Count - Threading impact
  12. Cache Misses vs Message Size - Memory hierarchy behavior
  13. CPU Cycles/Byte - Efficiency comparison

## [IMAGES: All 4 plots]





## Part E: Performance Analysis

### Question 1: Why doesn't zero-copy always give best throughput?

#### Answer:

Zero-copy underperforms on this system for these reasons:

#### 14. Loopback Testing Limitation

- Tests use network namespaces on same machine
- No real NIC DMA → kernel still copies data internally
- Page pinning overhead WITHOUT DMA benefit

#### 15. Page Pinning Cost (~2500 cycles per send)

- Walk page tables
- Increment refcount
- Lock pages in memory
- For small messages (<16KB), this overhead > memcpy cost

#### 16. Async Completion Overhead

- Must drain error queue (MSG\_ERRQUEUE)
- Adds latency: 65.54 us vs 19.22 us @ 64KB
- Extra context switches

#### Evidence:

Message Size	Two-Copy	Zero-Copy	Difference
1 KB	2.63 Gbps	0.91 Gbps	-65%
64 KB	43.90 Gbps	22.91 Gbps	-48%

**When zero-copy WOULD win:** Real NIC on 10GbE+, messages >64KB

---

---

## Question 2: Which cache level shows most reduction and why?

**Answer:**

**L1 cache** shows the most reduction (40% fewer misses for one-copy).

**Evidence @ 16KB, 4 threads:**

Implementation	L1 Misses	Reduction	LLC Misses	Change
Two-Copy	25,973,959	-	41,239,234	-
One-Copy	15,503,925	-40.3%	68,322,464	+65.7%

**Why L1 benefits:**

Two-Copy working set = Original (16KB) + Intermediate (16KB) = 32KB

One-Copy working set = Original (16KB) only = 16KB

L1 Cache size = 32KB

- Two-copy exceeds L1 capacity → thrashing
- One-copy fits in L1 → fewer misses

**Why LLC misses INCREASE (unexpected):**

- One-copy uses scattered memory (8 separate fields)
- Defeats hardware prefetcher (works best on sequential access)
- TLB pressure (8 pages vs 4 pages)
- Non-inclusive LLC can't help scattered evictions

**Lesson:** Memory access pattern matters more than copy count at higher cache levels.

---

---

### **Question 3: How does thread count interact with cache contention?**

**Answer:**

Thread count causes **super-linear performance degradation** due to cache contention.

#### **Evidence - Per-Thread L1 Misses @ 16KB:**

Threads	Total Misses	Per-Thread	Increase Factor
1	1,203,686	1,203,686	1×
4	25,973,959	6,493,490	5.4× (not 1×!)
8	468,128,002	58,516,000	48× (massive!)

#### **Mechanisms:**

##### 17. False Sharing - Threads write to nearby cache lines

```c

```
__sync_fetch_and_add(&total_bytes_received, bytes); // Cache line bouncing
```

...

- Thread 1 modifies → cache line in "Modified" state
- Thread 2 reads → invalidates Thread 1's cache line
- Constant cache coherence traffic

##### 18. Working Set Expansion

- 1 thread: 20KB fits in L1 (32KB)
- 4 threads: 80KB exceeds L1 → spills to LLC
- 8 threads: 160KB heavy LLC contention

## 19. Context Switching

- 1 thread: 203 context switches
- 8 threads: 487,281 switches ( $63\times$  more!)
- Each switch = TLB flush + cold cache

**Optimal thread count for this workload: 2-4 threads**

---

## Question 4: At what message size does one-copy outperform two-copy?

**Answer:**

**One-copy outperforms at 64KB and above.**

**Throughput @ 4 threads:**

| Size  | Two-Copy   | One-Copy   | Winner   |
|-------|------------|------------|----------|
| 1 KB  | 2.63 Gbps  | 2.33 Gbps  | Two-Copy |
| 4 KB  | 7.03 Gbps  | 6.01 Gbps  | Two-Copy |
| 16 KB | 22.31 Gbps | 19.98 Gbps | Two-Copy |
| 64 KB | 43.90 Gbps | 46.89 Gbps | One-Copy |

**Why one-copy LOSES at small sizes:**

- Scatter-gather overhead: Setup 8 iovec entries (128 bytes metadata)
- At 1KB: overhead ratio = 12.8%
- Kernel path for sendmsg() is more complex than send()
- 8 scattered regions defeat prefetcher

### **Why one-copy WINS at 64KB:**

- Overhead amortized: 128 bytes / 65536 bytes = 0.2%
- Eliminates memcpy (saves 50% of data movement)
- Memory bandwidth approaching saturation (187 Gbps)
- Copy elimination becomes critical

**Crossover point: ~32-48 KB** (extrapolated)

---

### **Question 5: At what message size does zero-copy outperform two-copy?**

**Answer:**

**Zero-copy NEVER outperforms two-copy on this system.**

### **Throughput @ 4 threads:**

| Size  | Two-Copy   | Zero-Copy  | Difference |
|-------|------------|------------|------------|
| 1 KB  | 2.63 Gbps  | 0.91 Gbps  | -65%       |
| 4 KB  | 7.03 Gbps  | 2.74 Gbps  | -61%       |
| 16 KB | 22.31 Gbps | 9.13 Gbps  | -59%       |
| 64 KB | 43.90 Gbps | 22.91 Gbps | -48%       |

Best case: Zero-copy reaches only 52% of two-copy throughput.

### **Root Cause:**

**Loopback interface** removes DMA benefit:

Real Network: User memory → DMA (no CPU copy) → NIC → Network

Loopback: User memory → Copy happens anyway (no real NIC) → Same machine

Result: Page pinning overhead + No DMA benefit = Worst of both worlds

#### **Cost breakdown:**

- Page pinning: ~2500 cycles
- Completion notification: ~1000 cycles
- Total overhead: ~3500 cycles
- But data still gets COPIED in loopback (no DMA)

#### **When zero-copy WOULD work:**

- Real 10GbE NIC (not loopback)
  - Two physical machines
  - Message size >16KB
  - Expected speedup: 1.5-2.5× for bulk transfers
- 

## **Question 6: Identify one unexpected result and explain it**

#### **Answer:**

#### **Unexpected Result:**

One-copy shows **65% MORE LLC cache misses** than two-copy (68M vs 41M @ 16KB), despite eliminating a memory copy.

**Expected:** Fewer copies → Fewer cache misses

**Observed:** L1 improved (-40%) but LLC degraded (+65%)

## **Explanation:**

### **Hardware Prefetcher Behavior:**

Two-copy (Sequential access):

[c]

```
memcpy(buffer, field1, 2KB); // Sequential addresses  
memcpy(buffer+2KB, field2, 2KB); // Still sequential  
send(buffer, 16KB); // One contiguous block
```

- Prefetcher detects stride=64B
- Prefetches 8-16 lines ahead into LLC
- High LLC hit rate

One-copy (Scattered access):

[c]

```
iov[0].base = field1; // Address: 0x7f0000  
iov[1].base = field2; // Address: 0x7f2000 (8KB apart!)  
...  
sendmsg(iov[8]); // Kernel reads from 8 scattered regions
```

- Irregular stride confuses prefetcher
- No prefetching to LLC
- Many LLC misses

## **Additional Factors:**

### 20. TLB Pressure

- Two-copy: 4 pages (16KB / 4KB per page)
- One-copy: 8+ pages (8 separate malloc'd fields)
- TLB miss → 4-level page table walk → LLC traffic

## 21. Cache Line Alignment

- Two-copy: malloc() returns cache-aligned buffer (64B boundary)
- One-copy: 8 separate malloc() may return unaligned addresses
- Partial cache lines reduce effective LLC capacity

## 22. Non-Inclusive LLC (modern Intel)

- LLC is "victim cache" for L1/L2 evictions
- Sequential access plays nice with victim cache
- Scattered access creates more evictions → LLC overflow

### Conclusion:

Modern CPUs optimize for sequential access patterns (prefetching, cache lines, TLB). "Optimizations" that scatter data can backfire at higher cache levels. Always profile multiple cache levels!

---

## AI Usage Declaration

### Tools Used

GitHub Copilot - AI coding assistant

### What AI Helped With

| Component    | AI Assistance                                        | Student Work                                  |
|--------------|------------------------------------------------------|-----------------------------------------------|
| C Programs   | Socket boilerplate, error handling, pthread patterns | Logic, architecture, performance optimization |
| Bash Script  | Script structure, perf parsing                       | Network namespaces, experimental design       |
| Python Plots | Matplotlib syntax, formatting                        | Data collection, analysis,                    |

|        |                               |                                      |
|--------|-------------------------------|--------------------------------------|
|        |                               | visualization choices                |
| Report | Markdown formatting, diagrams | ALL analysis, explanations, insights |

### Specific Prompts Used:

23. "Create TCP server socket with pthread threading"
  24. "Implement sendmsg() with iovec for scatter-gather"
  25. "Implement MSG\_ZEROCOPY with completion handling"
  26. "Parse perf stat output to extract CPU cycles and cache misses"
  27. "Create plotting functions for throughput vs message size"
- 
- 

## Key Findings Summary

### Performance Winners by Message Size

| Size      | Best Approach | Throughput | Reason                     |
|-----------|---------------|------------|----------------------------|
| 1-16 KB   | Two-Copy      | 22.31 Gbps | Simple, cache-friendly     |
| 64 KB+    | One-Copy      | 46.89 Gbps | Eliminates memcpy overhead |
| All sizes | ~~Zero-Copy~~ | 22.91 Gbps | Loopback limitation        |

### Top Insights

28. Cache pattern > Copy count - Sequential access beats fewer copies at LLC level
29. Threading sweet spot: 2-4 threads - Beyond this, cache contention dominates
30. L1 vs LLC trade-off - One-copy improves L1 (-40%) but hurts LLC (+65%)
31. Zero-copy needs real hardware - Loopback testing defeats the purpose
32. Hardware prefetcher is critical - Scattered access kills performance

### Real-World Recommendations

- Small packets (<16KB): Use two-copy (simple, fast)
- Large transfers (>64KB): Use one-copy on loopback, zero-copy on real NICs
- Thread count: Match to core count, don't exceed it
- Always profile multiple cache levels: L1 gains may not translate to overall wins
-

**GitHub Repository:** [https://github.com/shivam697/GRS\\_PA02](https://github.com/shivam697/GRS_PA02)

---

*End of Report*